

BRIAN REVIEW / JULY 21, 2026

PERFORMANCE PROOF PACK

Legacy baseline vs. the Sinister Diesel V2 release candidate — measured, attributed, and ready for executive review.

V2

MEASURE / VERIFY / RELEASE

90%

GTMETRIX STRUCTURE

52.8%

PAYLOAD REDUCTION

51.5%

REQUEST REDUCTION

TUNE

CURRENT VERDICT

Release evidence

Prepared from direct V2 preview measurement, the supplied legacy baseline, source-level regression evidence, and clearly stated limitations.

01 Executive Verdict

The evidence supports a precise conclusion—not a blanket speed claim.

Decision statement

V2 is materially leaner, structurally stronger, and more visually stable. It is not yet faster on initial render: Performance and LCP still require an optimization pass before publication can claim an overall speed win.

77 to 90

STRUCTURE

618 to 558ms

BLOCKING TIME


.02 to .01

LAYOUT SHIFT

59 to 51

PERFORMANCE

Why this is still meaningful: V2 cuts the median transferred payload from 5.29 MB to 2.50 MB and median requests from 169 to 82 under identical local conditions. Those are foundational engineering gains. The remaining delay is concentrated in server response and the critical render path, making the next optimization phase specific and measurable.




Directly captured V2 preview homepage at 1440 × 900. Preview credential omitted from all artifacts.

02 GTmetrix Head-to-Head

Two remote tests, fifteen minutes apart, from Seattle using Chrome 142 and Lighthouse 12.6.1.

Metric	Legacy 3:07 PM	V2 3:22 PM	Delta	Verdict
GTmetrix grade	D	D	—	UNCHANGED
Performance	59%	51%	−8 pts	NEEDS WORK
Structure	77%	90%	+13 pts	V2 WIN
Largest Contentful Paint	2.1s	2.2s	+0.1s	ESSENTIALLY FLAT
Total Blocking Time	618ms	558ms	−60ms	V2 WIN
Cumulative Layout Shift	0.02	0.01	−0.01	V2 WIN




What GTmetrix proves

V2 improves page construction quality, blocking time, and visual stability. It also proves the initial-render path is not yet optimized enough to outperform legacy on the composite Performance score.

03 Controlled Six-Run Comparison

Three fresh desktop runs per version under identical local Lighthouse 13.4.1 conditions; medians shown.

Metric	Legacy median	V2 median	Change	Interpretation
Performance score	93	85	−8 pts	Legacy faster initially
First Contentful Paint	1.09s	1.60s	+46.4%	Optimize V2 render path
Largest Contentful Paint	1.39s	1.60s	+14.9%	Optimize V2 LCP
Speed Index	1.37s	2.17s	+58.3%	Visual completion slower
Cumulative Layout Shift	0.034	0.008	−76.5%	Major V2 stability win
Transferred bytes	5.29 MB	2.50 MB	−52.8%	Major V2 payload win
Network requests	169	82	−51.5%	Major V2 request win
Server response	883ms	907ms	+2.7%	Platform/server opportunity



The local score differs from GTmetrix because Lighthouse version, execution environment, and remote network/server conditions differ. The fair comparison is within each environment: GTmetrix-to-GTmetrix and local-median-to-local-median, never one tool's number against the other tool's number.

04 Why V2 Is Better Engineered

The strongest V2 gains are structural and operational, not yet an across-the-board speed win.


Evidence	Measured change	Customer/business significance
GTmetrix Structure	77% → 90%	Cleaner technical construction and fewer structural penalties.
Transferred payload	5.29 MB → 2.50 MB median	Less data for customers to download; more headroom for mobile and slower connections.
Network requests	169 → 82 median	Fewer connection, queueing, and third-party failure opportunities.
Layout stability	CLS .034 → .008 local; .02 → .01 remote	Less visual movement while customers read and interact.
Remote blocking time	618ms → 558ms	Less time where the main thread cannot respond.
Commerce regression coverage	13+ source contracts plus route audits	Changes are safer to release and diagnose than a screenshot-only redesign.
SEO controls	Dynamic metadata, canonicals, robots, schema	A stronger technical foundation for production crawling and indexing.

Credible positioning for Brian

V2 is the stronger platform and the better customer experience. Before promising that it is the faster homepage, complete the focused performance phase and retest against this baseline.

05 SEMrush Baseline and V2 SEO Readiness

The supplied dashboard describes the production legacy domain. A private preview cannot have independent rankings or traffic.



Legacy SEMrush baseline	Observed value	V2 position before publish
Authority Score	33	Domain-level value; shared until V2 becomes production
Organic keywords	11.8K	Cannot be attributed to private preview
Referring domains	1.1K	Domain asset; preserved through URL/canonical discipline
Site Health	83% across 100 pages	Re-crawl after production publication
Errors / warnings	26 / 1,332	Resolve from post-launch V2 crawl, not estimated
Position visibility	52.94%	Measure after indexing stabilizes
AI visibility	37	Production-domain outcome; monitor post-launch

V2 source-verified SEO controls

- Dynamic titles and meta descriptions with safe fallbacks
- Canonical URLs for indexable commerce and content routes
- noindex,follow for private, account, basket, checkout and internal-search routes
- Product, category, breadcrumb and organization structured-data paths
- Button/link integrity and exactly-one-H1 regression contracts
- Encoding and editorial-route regression coverage

What cannot be claimed yet

No responsible report can claim higher V2 rankings, authority, traffic, conversions, or SEMrush Site Health before V2 is published and crawled. This pack records readiness, not invented outcomes.

06 Focused Performance Roadmap

The audit isolates a practical sequence for converting structural gains into a speed win.

Priority	Action	Evidence	Expected direction
P0	Reduce initial HTML/server response latency	910 ms median-run observation	Earlier FCP/LCP
P0	Split or defer noncritical JavaScript	Estimated savings 474 KiB	Lower parse/evaluation and blocking
P0	Extract critical CSS and load page-family CSS intentionally	Estimated savings 154 KiB	Faster first render
P1	Remove render-blocking requests	Estimated savings 380 ms	Potential ~380ms lab saving
P1	Minify CSS and JavaScript	37 KiB CSS + 10 KiB JS opportunity	Smaller transfer/parse cost
P1	Optimize hero/LCP image delivery	Estimated savings 95 KiB	Faster LCP without redesign
P2	Retest GTmetrix three times after changes	Use same Seattle/Chrome class	Prove the speed claim

Acceptance target

Retain Structure ≥ 90 , CLS ≤ 0.01 and the ~50% payload/request reductions while raising GTmetrix Performance above 59 and reducing LCP below 2.1s under the same test class.

07 Methodology, Limits, and Audit Trail

Evidence class	Source	Use in this report
Legacy GTmetrix	User-supplied report generated 3:07 PM	Baseline remote performance
V2 GTmetrix	Direct report q7RFhvs generated 3:22 PM	Remote V2 performance; credential redacted
Controlled Lighthouse	3 V2 + 3 legacy fresh desktop runs	Median payload, requests, paint and stability comparison
Legacy SEMrush	User-supplied July 21 dashboard	Production domain baseline only
V2 SEO readiness	Repository tests and template inspection	Implemented controls, not ranking claims
Visual evidence	Direct V2 preview capture	Confirms tested release-candidate identity

GTmetrix remote tests are highly comparable because they used the same location, browser version, Lighthouse version, target domain, and a fifteen-minute window. The V2 report was run through the preview state and its credential was removed from saved evidence. Local Lighthouse medians are a second evidence track and are never mixed numerically with GTmetrix scores.

Performance varies with origin response, third-party availability, cache state, and test infrastructure. A single test is evidence, not a permanent guarantee. Production Core Web Vitals require real traffic after publication.

Final recommendation

Approve V2 as the stronger storefront platform and customer experience, subject to the documented release gates.
Complete the focused performance optimization pass before using 'faster' as a public or executive claim.